

# Python Exercises

## Zenoh - OpenCV - MCAP

.....

This worksheet introduces three tools for building distributed and multimedia applications in Python:

- **Zenoh** – A communication middleware that unifies data in motion, data at rest, and computations. It supports publish/subscribe, queries, and storage in an efficient and flexible way.
- **OpenCV** – A widely used computer vision library for image and video processing, machine learning, and AI applications.
- **MCAP** – A flexible, performant, and well-structured container format for recording and replaying message data (similar to ROS bag files but more general-purpose).

The goal of these exercises is to gain hands-on experience in integrating these technologies. By the end, you will be able to:

- Publish and subscribe to messages using Zenoh.
- Capture and process video frames with OpenCV.
- Record and replay data using MCAP.

Resources:

- **Zenoh**: <https://zenoh.io>
- **OpenCV**: <https://docs.opencv.org/4.12.0/>
- **MCAP**: <https://mcap.dev/docs/python/>

## 1. Quick Start: Structuring Your Python Project

To keep your code organized, use the following folder structure:

```
zenoh_opencv_mcap/  
├── README.md  
├── requirements.txt  
├── venv/                # Python virtual environment  
├── src/  
│   ├── pub_sub.py      # Source files goes here  
│   └── ...
```

- `src/` will hold your source code.
- `requirements.txt` will specify dependencies.
- `venv/` is the isolated Python environment (so you don't pollute system packages).

### 1.1. Create and Activate a Virtual Environment

```
1 # Create a project folder  
2 mkdir zenoh_opencv_mcap  
3 cd zenoh_opencv_mcap  
4  
5 # Create a virtual environment  
6 python3 -m venv venv  
7  
8 # Activate it  
9 source venv/bin/activate
```

bash

When activated, your terminal prompt should show (venv) at the start.

## 1.2. Install Required Packages

We need Zenoh Python API, OpenCV, and MCAP.

```
1 pip install eclipse-zenoh opencv-python mcap
```

bash

Optionally, create a requirements.txt file:

```
1 pyzenoh
2 opencv-python
3 mcap
```

txt

Then install all with:

```
1 pip install -r requirements.txt
```

bash

## 1.3. Verify Installations

In Python:

```
1 import zenoh
2 import cv2
3 import mcap
4 print("Zenoh module loaded successfully:")
5 print("OpenCV version:", cv2.__version__)
6 print("MCAP version:", mcap.__version__)
```

python

If no errors occur, you are ready to start the exercises.

## 2. Zenoh Basics

Zenoh provides a unified way to handle publish/subscribe (pub/sub) and query/reply. These exercises will help you learn the fundamentals of using Zenoh in Python.

### 2.1. First Publisher

Write a Python program that publishes a message every second. Create a file src/publisher.py:

```
1 import time
2 import zenoh
3
4 def main():
5     z = zenoh.open(zenoh.Config()) # Open a Zenoh session
6     pub = z.declare_publisher("demo/example/hello") # Declare a publisher
7     for i in range(10): # Publish messages every second
8         msg = f"Hello {i}"
9         print("[Publisher] Publishing:", msg)
10        pub.put(msg)
11        time.sleep(1)
12    z.close() # Close session
13 if __name__ == "__main__":
14    main()
```

python

Run it in one terminal:

```
1 python src/publisher.py
```

bash

## 2.2. First Subscriber

```
1 import time
2 import zenoh
3
4 def main():
5     z = zenoh.open(zenoh.Config()) # Open a Zenoh session
6     # Declare a subscriber
7     def listener(sample):
8         print(f"[Subscriber] Received: {sample.payload.to_string()}")
9
10    sub = z.declare_subscriber("demo/example/hello", listener)
11    print("Listening... Press Ctrl+C to stop.")
12    try:
13        while True:
14            time.sleep(1)
15    except KeyboardInterrupt:
16        pass
17    z.close() # Close session
18
19 if __name__ == "__main__":
20     main()
```

python

Run it in another terminal at the same time as the publisher. You should see the subscriber printing messages.

## 2.3. Request/Reply

Zenoh also supports a query/reply model.

Create a file `src/responder.py`:

```
1 import time
2 import zenoh
3
4 def main():
5     z = zenoh.open(zenoh.Config())
6     # Register a queryable resource
7     def callback(query):
8         print("[Responder] Received query:", query.selector)
9         query.reply(query.selector.key_expr, "This is PSR")
10
11    qable = z.declare_queryable("demo/example/service", callback)
12    print("Responder ready. Ctrl+C to stop.")
13    try:
14        while True:
```

python

```

15         time.sleep(1)
16     except KeyboardInterrupt:
17         pass
18
19 cv2.destroyAllWindows()
20     z.close()
21
22 if __name__ == "__main__":
23     main()

```

src/requester.py

```

1  import zenoh
2
3  def main():
4      z = zenoh.open(zenoh.Config())
5
6      q = z.declare_querier("demo/example/service")
7      replies = q.get()
8      for reply in replies:
9          print("[Requester] Got reply:", reply.ok.payload.to_string())
10
11     z.close()
12
13 if __name__ == "__main__":
14     main()

```

python

## 2.4. Challenges

### 2.4.1. Countdown Publisher

Write a new program called `countdown_publisher.py` that:

- Publishes the numbers from 10 down to 1.
- Waits one second between each publish.
- Publishes the final message “Lift-off!”.

Run it together with your subscriber from Exercise 2.

### 2.4.2. Math Service

Implement a Zenoh service called “demo/math/square” that:

- Waits for a query containing a number.
- Replies with the square of that number.
- For example, a query with “7” should reply “49”.

Run your `math_responder.py` in one terminal, and write a `math_requester.py` that asks for numbers 1 through 5 and prints the replies.

## 3. Combining Zenoh and OpenCV

In these exercises, you will learn how to publish video frames captured with OpenCV and subscribe to them using Zenoh.

### 3.1. Publish Camera Frames

Create a new file `src/video_publisher.py`:

```
1  import cv2
2  import numpy as np
3  import zenoh
4  import time
5
6  def main():
7      z = zenoh.open(zenoh.Config())
8
9      def listener(sample):
10         # Convert received bytes back into an image
11         np_arr = np.frombuffer(sample.payload.to_bytes(), np.uint8)
12         frame = cv2.imdecode(np_arr, cv2.IMREAD_COLOR)
13
14         # Show frame
15         cv2.imshow("Subscriber Video", frame)
16         if cv2.waitKey(1) & 0xFF == ord("q"):
17             exit(0)
18
19     sub = z.declare_subscriber("demo/video/frame", listener)
20     print("Listening for frames... Press 'q' to quit")
21     try:
22         while True:
23             time.sleep(1)
24     except KeyboardInterrupt:
25         pass
26     cv2.destroyAllWindows()
27     z.close()
28
29 if __name__ == "__main__":
30     main()
```

Run it, and it will continuously capture frames and publish them as JPEG-encoded bytes.

### 3.2. Subscribe and Display Frames

Create `src/video_subscriber.py`:

```
1  import cv2
2  import zenoh
3  import time
4
5  def main():
6      z = zenoh.open(zenoh.Config())
7      pub = z.declare_publisher("demo/video/frame")
8      cap = cv2.VideoCapture(0) # Open default camera (0 = webcam)
```

```

9     if not cap.isOpened():
10         print("X Cannot open camera")
11         return
12     while True:
13         ret, frame = cap.read()
14         if not ret:
15             print("X Failed to grab frame")
16             break
17         _, buffer = cv2.imencode(".jpg", frame) # Encode frame as JPEG
18         pub.put(buffer.tobytes()) # Publish frame
19         print("[Publisher] Sent frame")
20     cap.release()
21     z.close()
22
23 if __name__ == "__main__":
24     main()

```

Run this in a separate terminal (or even on a different machine in the same network). You should see the video feed from the publisher.

### 3.3. Apply OpenCV Processing Before Publishing

Modify the publisher to grayscale the frames before sending.

In `src/video_publisher.py`, inside the loop:

```

1 # Convert to grayscale
2 gray = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)
3
4 # Encode as JPEG
5 _, buffer = cv2.imencode(".jpg", gray)
6
7 # Publish
8 pub.put(buffer.tobytes())
9 print("[Publisher] Sent grayscale frame")

```

python

### 3.4. Face Detection with OpenCV + Zenoh

Let's process frames with a Haar cascade classifier for face detection.

Update publisher loop:

```

1
2 face_cascade = cv2.CascadeClassifier(cv2.data.harcascades +
3     "haarcascade_frontalface_default.xml")
4
5 while True:
6     ret, frame = cap.read()
7     if not ret:
8         break

```

python

```

9     faces = face_cascade.detectMultiScale(gray, 1.3, 5)
10    # Draw rectangles around detected faces
11    for (x, y, w, h) in faces:
12        cv2.rectangle(frame, (x, y), (x+w, y+h), (255, 0, 0), 2)
13
14    _, buffer = cv2.imencode(".jpg", frame)
15    pub.put(buffer.tobytes())
16    print(f"[Publisher] Sent frame with {len(faces)} faces detected")

```

The subscriber will now show a live video stream with bounding boxes around faces.

## 3.5. Challenges

### 3.5.1. Edge Detection

Update your video publisher to apply an edge detection filter (Canny) before publishing. The subscriber should display a live video showing only the detected edges.

### 3.5.2. Moving Circle

Draw a red circle that moves horizontally across the video (like an animation) and resets when it reaches the edge. Overlay it on top of the live camera feed before publishing.

## 4. Recording and Replaying with MCAP

### 4.1. Record Frames into MCAP

We'll subscribe to the Zenoh video stream and write each received frame into an MCAP file.

Create `src/video_recorder.py`:

```

1  import zenoh
2  import time
3  from mcap.writer import Writer
4
5  def main():
6      z = zenoh.open(zenoh.Config())
7      # Open MCAP file for writing
8      with open("video_stream.mcap", "wb") as f:
9          writer = Writer(f)
10         start_time = int(time.time() * 1e9) # nanoseconds
11
12         # Create channel for video frames
13         channel_id = writer.register_channel(
14             topic="demo/video/frame",
15             message_encoding="opencv_jpg",
16             schema_id=0,
17         )
18
19         def listener(sample):
20             ts = int(time.time() * 1e9) # timestamp in ns
21             data = sample.payload.to_bytes()

```

python

```

22         writer.add_message(
23             channel_id=channel_id,
24             log_time=ts,
25             publish_time=ts,
26             data=data,
27         )
28         print(f"[Recorder] Wrote frame at {ts}")
29
30     writer.start()
31     sub = z.subscribe("demo/video/frame", listener)
32
33     print("Recording... Press Ctrl+C to stop")
34     try:
35         while True:
36             time.sleep(1)
37     except KeyboardInterrupt:
38         pass
39
40     # Finish writing MCAP file
41     writer.finish()
42     print("Saved to video_stream.mcap")
43
44 if __name__ == "__main__":
45     main()

```

Run the publisher, the subscriber, and then run video\_recorder.py. It will save all frames to video\_stream.mcap.

## 4.2. Replay Frames from MCAP

Now, we'll replay the recorded frames from the MCAP file and display them.

Create src/video\_player.py

```

1  import cv2
2  import numpy as np
3  from mcap.reader import make_reader
4
5  def main():
6      with open("video_stream.mcap", "rb") as f:
7          reader = make_reader(f)
8
9          for schema, channel, message in reader.iter_messages():
10             frame = cv2.imdecode(
11                 np.frombuffer(message.data, np.uint8), cv2.IMREAD_COLOR
12             )
13             cv2.imshow("MCAP Replay", frame)
14

```



```
15         # Wait ~33 ms for ~30 FPS replay
16         if cv2.waitKey(33) & 0xFF == ord("q"):
17             break
18
19     cv2.destroyAllWindows()
20
21 if __name__ == "__main__":
22     main()
```

Run `video_player.py` and you should see the replay of the previously recorded video.

## 4.3. Challenges

### 4.3.1. Replay into Zenoh

Instead of just displaying frames locally, let's replay them back into Zenoh so that other subscribers (e.g., `video_subscriber.py`) can see them as if they were live.

### 4.3.2. Adjustable Replay Speed

Modify your `video_replayer.py` so that it accepts a command-line argument for speed multiplier:

- 1.0 → real-time speed
- 0.5 → half speed (slower)
- 2.0 → double speed (faster)

Test with your previously recorded file.